

# First steps in MPI

Joachim Hein (LUNARC, Lund University)

Xin Li (PDC, KTH)

## Overview

- This lecture: Basic set up inside the code for MPI
- Header files
- Initialisation of the MPI library
- Finalisation of the MPI library

## Header files

- Every compilation module accessing MPI requires inclusion of a header file:

- F77 style:

```
include "mpif.h"
```

- Fortran90:

```
use mpi
```

- Fortran08, **MPI 3.0**, Fortran standard compliant!

```
use mpi_f08
```

- C:

```
#include "mpi.h"
```

## mpi4py

- Python code does not require compilation or header file, but the MPI module needs to be imported

- Python:

```
from mpi4py import MPI
```

## MPI command in C

- In C all MPI commands are functions with return type **int**

```
int MPI_Abcddef( arguments )
```

- The returned value is the error code
  - Detailing problems with the command
- Typically very hard to recover from MPI-errors
- Most codes do not check these error codes
- **Rem:** MPI commands can modify arguments
  - pass a pointer

## MPI command in Fortran

- In Fortran all MPI commands are subroutines

```
MPI_ABCDEF( arguments, ierror )
```

- MPI commands in Fortran carry **one more argument** than their C counter part
  - This is optional in Fortran 2008
  - This is of type **int** and returns the error code
- Again, this is typically unchecked, hence easily forgotten while coding
- Forgetting this in F77/F90 typically leads to segmentation faults at runtime

## MPI command in Python

- In Python all MPI commands are methods of MPI communicator

`comm.method(arguments)`

- In general fewer arguments are needed, compared to C and Fortran
- Communication of generic Python objects is done via all-lowercase methods (e.g. `comm.send(...)` )
- Communication of buffer-like objects is done via methods with an uppercase letter (e.g. `comm.Send(...)` )

## C++ bindings: depreciated/removed

- MPI used to have special C++ bindings
- Depreciated since MPI standard 2.2                      September 2009
- Removed in MPI standard 3.0                              September 2012
- Use C bindings in C++ programs
  - Consider wrapping in OO-style for your app's needs

## MPI\_Init

- The first MPI call of any MPI program has to be **MPI\_Init**
- In C:

```
int MPI_Init(int *argc, char ***argv)
```

- Arguments are same as **main**
- Alternatively modern MPI libraries allow to pass **Null**

- In Fortran

```
MPI_INIT(IERROR)
INTEGER:: IERROR
```

## MPI\_Finalize

- The last MPI call has to be **MPI\_Finalize**
- In C:

```
int MPI_Finalize(void)
```

- In Fortran:

```
MPI_FINALIZE(IERROR)
INTEGER:: IERROR
```

## MPI init and finalize in Python

- In Python, MPI is initialized upon import, and finalized upon exit

```
from mpi4py import MPI
```

## Minimal program in C

```
#include "mpi.h"

int main(int argc, char **argv)
{
    MPI_Init(&argc, &argv); // alt.: NULL,NULL

    // further MPI calls go here!

    MPI_Finalize();
}
```

## Minimal program in Fortran

```
program minimpi
  use mpi      ! alt.: include "mpif.h"
  integer:: ierror

  call MPI_INIT(ierror)
           ! further MPI calls go here
  call MPI_FINALIZE(ierror)

end program minimpi
```

## Minimal program in Python

```
from mpi4py import MPI

# further MPI calls go here
```

## Summary

- Basic requirements for an MPI program
  - Header files
  - Initialising MPI
  - Finalising MPI

## Communicators



## Overview


- Concept of communicators
- Predefined communicator
- Querying basic properties of the communicator

## Communicator

- Most messages passed inside (intra-)communicator
- Communicator
  - Group of processes
  - Save communication universe
  - Order
  - Can have additional topology


## Example communicator

- Picture shows:
  - 9 processes
  - 3 communicators
- Processes carry label
  - Here: labels for red communicator
- Labels start at 0




## Example communicator

- Picture shows:
  - 9 processes
  - 3 communicators
- Processes carry label
  - Here: labels for violet communicator
- Labels start at 0




## Example communicator

- Picture shows:
  - 9 processes
  - 3 communicators
- Processes carry label
  - Here: labels for blue communicator
- Labels start at 0
- **Label depends on the communicator**



## Predefined communicator **MPI\_COMM\_WORLD**

- After **MPI\_Init** one predefined communicator:  
**MPI\_COMM\_WORLD**  
 (in Python: **MPI.COMM\_WORLD**)
- This contains all processes
- In C, this is typedef: **MPI\_Comm**
- In Fortran90 this is: **INTEGER**
- In Fortran 2008 this is: **type (MPI\_Comm)**
- Further communicators: user created



## MPI\_Comm\_size

- Number of processes in a communicator
- In C:

```
int MPI_Comm_size(MPI_Comm comm, int *size)
```

- In Fortran 90:

```
MPI_COMM_SIZE(COMM, SIZE, IERROR)  
INTEGER:: COMM, SIZE, IERROR)
```

- Arguments:
  - comm**: communicator (input)
  - size**: number of processes (output)

## Get\_size in Python

- Number of processes in a communicator

```
comm.Get_size()
```

- No arguments
- Returns the number of processes in a communicator

## MPI\_Comm\_rank

- Rank (label) of the process
- In C:

```
int MPI_Comm_rank(MPI_Comm comm, int *rank)
```

- In Fortran 90:

```
MPI_COMM_RANK(COMM, RANK, IERROR)  
INTEGER :: COMM, RANK, IERROR)
```

- Arguments:
  - comm**: communicator (input)
  - rank**: rank of processes (output)

## Get\_rank in Python

- Rank (label) of the process

```
comm.Get_rank()
```

- No arguments
- Return the rank of this process in a communicator

## Copying communicators

- Extensive use of **MPI\_COMM\_WORLD** is *discouraged*
- Exactly **one** reference to **MPI\_COMM\_WORLD** in the program (apart from **MPI\_Abort**):

- Copy it:

```
my_world = MPI_COMM_WORLD
```

- Use **my\_world** in the rest of the program
- Declare **my\_world** as
  - **MPI\_Comm** in C
  - **INTEGER** in Fortran 90
  - **type (MPI\_Comm)** in Fortran 08

## Copying communicators in Python

- Extensive use of **MPI.COMM\_WORLD** is *discouraged*
- Exactly **one** reference to **MPI.COMM\_WORLD** in the program (apart from **Abort**):

```
my_comm = MPI.COMM_WORLD
```

- Use **my\_comm** in the rest of the program

## MPI\_Abort

- Aborting all MPI tasks from any task (e.g. read corrupt input file, failed safety check)

```
int MPI_Abort(MPI_Comm comm, int errorcode)
```

In Fortran 90:

```
MPI_ABORT(COMM, ERRORCODE, IERROR)
```

```
INTEGER:: COMM, ERRORCODE, IERROR
```

- COMM is the communicator with the task to abort
  - typically `MPI_COMM_WORLD`
- ERRORCODE returned to the UNIX shell to flag a problem
- All arguments: input

## Abort in Python

- Aborting all MPI tasks from any task (e.g. read corrupt input file, failed safety check)

```
comm.Abort(errorcode)
```

- Typically called by `MPI.COMM_WORLD`
- errorcode returned to the UNIX shell to flag a problem

## Summary

- Concept of communicator
- Predefined communicator `MPI_COMM_WORLD` (`MPI.COMM_WORLD` in Python)
- Querying task rank and size of a communicator
- Aborting a program on error
- You should now be able to write simple MPI programs, which are useful (e.g. task farm)