

Introduction to Point-to-Point Communication

Joachim Hein (LUNARC, Lund University)

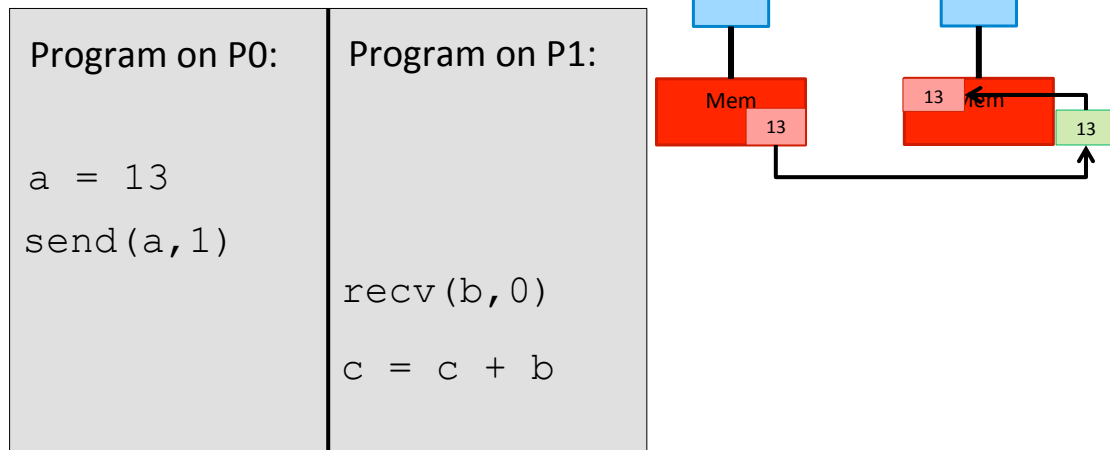
Xin Li (PDC, KTH)

Overview

- Concept: Message passing point-to-point
- The MPI application interface for sending and receiving messages

Double sided point-to-point communication

- Most basic form of communication in message passing:



Standard send in C: MPI_Send

```
int MPI_Send(void* buf, int count, MPI_Datatype
datatype, int dest, int tag, MPI_Comm comm)
```

- **buf:** address of send buffer
- **count:** number of elements to be sent
- **datatype:** data type of buffer (explained further down)
- **dest:** rank of receiver
- **tag:** message tag (put 0 if you don't need)
- **comm:** communicator

Standard send Fortran 90: MPI_SEND

```
MPI_SEND(BUF, COUNT, DATATYPE, DEST, TAG, &
          COMM, IERROR)
```

```
<TYPE>:: BUF
```

```
INTEGER:: COUNT, DATATYPE, DEST, TAG, COMM, IERROR
```

- **BUF:** (address of) send buffer
- **COUNT:** number of elements to be sent
- **DATATYPE:** data type of buffer (explained further down)
- **DEST:** rank of receiver
- **TAG:** message tag (put 0 if you don't need)
- **COMM:** communicator

Standard send Fortran 2008: MPI_SEND

```
MPI_SEND(BUF, COUNT, DATATYPE, DEST, TAG, &
          COMM, IERROR)
```

```
TYPE(*), DIMENSION(..), INTENT(IN):: BUF
```

```
INTEGER, INTENT(IN):: COUNT, DEST, TAG
```

```
TYPE(MPI_Datatype), INTENT(IN):: DATATYPE
```

```
TYPE(MPI_Comm), INTENT(IN):: COMM
```

```
INTEGER, INTENT(OUT), OPTIONAL:: IERROR
```

- **BUF:** (address of) send buffer
- **COUNT:** number of elements to be sent
- **DATATYPE:** data type of buffer (explained further down)
- **DEST:** rank of receiver
- **TAG:** message tag (put 0 if you don't need)
- **COMM:** communicator

Standard send in Python: send

```
comm.send(obj, dest=dest, tag=tag)
```

- **obj:** The Python object being sent
- **dest:** The rank number of the rank the data will be sent to
- **tag:** Message tag (optional)

Example:

```
comm.send(obj, dest=1, tag=0)
```

Receiving data in C

```
int MPI_Recv(void* buf, int count, MPI_Datatype  
datatype, int source, int tag, MPI_Comm comm,  
MPI_Status *status)
```

- **buf:** address of receive buffer ([output](#))
- **count:** number of elements to be received
- **datatype:** data type of buffer (explained further down)
- **source:** rank of sender (data origin)
- **tag:** message tag (needs to match the send!)
- **comm:** communicator
- **status:** status ([output](#)), info on: sender, tag, error

Receiving data in Fortran 90

```
MPI_RECV(BUF, COUNT, DATATYPE, SOURCE, TAG, &
          COMM, STATUS, IERROR)
```

```
<TYPE>:: BUF
```

```
INTEGER:: COUNT, DATATYPE, SOURCE, TAG, COMM,
          STATUS(MPI_STATUS_SIZE), IERROR
```

- BUF: address of receive buffer ([output](#))
- COUNT: number of elements to be received
- DATATYPE: data type of buffer (explained further down)
- SOURCE: rank of sender (data origin)
- TAG: message tag (needs to match the send!)
- COMM: communicator
- STATUS: status ([output](#)), info on: sender, tag, error

Receiving data in Fortran 2008

```
MPI_RECV(BUF, COUNT, DATATYPE, SOURCE, TAG, &
          COMM, STATUS, IERROR)
```

```
TYPE(*), DIMENSION(..) :: BUF
```

```
INTEGER, INTENT(IN) :: COUNT, SOURCE, TAG
```

```
TYPE(MPI_Datatype), INTENT(IN) :: DATATYPE
```

```
TYPE(MPI_Comm), INTENT(IN) :: COMM
```

```
TYPE(MPI_Status) :: STATUS
```

```
INTEGER, OPTIONAL, INTENT(OUT) :: IERROR
```

- BUF: address of receive buffer ([output](#))
- COUNT: number of elements to be received
- DATATYPE: data type of buffer (explained further down)
- SOURCE: rank of sender (data origin)
- TAG: message tag (needs to match the send!)
- COMM: communicator
- STATUS: status ([output](#)), info on: sender, tag, error

Receiving data in Python

```
obj = comm.recv(source=source, tag=tag)
```

- **obj**: The Python object being received
- **source**: rank of sender (data origin)
- **tag**: message tag (optional; needs to match the send)

Example:

```
myobj = comm.recv(source=0, tag=0)
```

Predefined data types in C (selection)

MPI datatype	C datatype
MPI_CHAR	signed char
MPI_SHORT	signed short int
MPI_INT	signed int
MPI_LONG	signed long int
MPI_FLOAT	float
MPI_DOUBLE	double
MPI_LONG_DOUBLE	long double

- choose the MPI datatype matching the send/receive buffer

Predefined data types in Fortran (selection)

MPI datatype	Fortran datatype
<code>MPI_INTEGER</code>	<code>INTEGER</code>
<code>MPI_REAL</code>	<code>REAL</code>
<code>MPI_DOUBLE_PRECISION</code>	<code>DOUBLE PRECISION</code>
<code>MPI_COMPLEX</code>	<code>COMPLEX</code>
<code>MPI_DOUBLE_COMPLEX</code>	<code>DOUBLE COMPLEX</code>
<code>MPI_LOGICAL</code>	<code>LOGICAL</code>

- choose the MPI datatype matching the send/receive buffer

Predefined data types in Python (selection)

mpi4py datatype	MPI datatype
<code>MPI.CHAR</code>	<code>MPI_CHAR</code>
<code>MPI.SHORT</code>	<code>MPI_SHORT</code>
<code>MPI.INT</code>	<code>MPI_INT</code>
<code>MPI.LONG</code>	<code>MPI_LONG</code>
<code>MPI.FLOAT</code>	<code>MPI_FLOAT</code>
<code>MPI.DOUBLE</code>	<code>MPI_DOUBLE</code>
<code>MPI.LONG_DOUBLE</code>	<code>MPI_LONG_DOUBLE</code>

- no need to specify the datatype for sending/receiving standard Python objects

The argument count (Fortran 77 style example)

- sample code:


```
integer :: a(9)
...
call mpi_send(a(3), 5, MPI_INTEGER, &
              2, 0, my_comm, merror)
```

- Where will it send?

- Sends to rank 2 of `my_comm`

- Which elements will it send?

- Sends elements: `a(3)`, `a(4)`, `a(5)`, `a(6)`, `a(7)`



Fortran kind syntax (F90 syntax) Real data

```
MPI_TYPE_CREATE_F90_REAL(P, R, NEWTYPE, IERROR)
```

```
INTEGER P, R, NEWTYPE, IERROR
```

- P: precision decimal digits
- R: exponent range
- NEWTYPE: requested MPI datatype (handle)
- Use to send real data declared with `selected_real_kind`
- Either P or R may be `MPI_UNDEFINED`

Example for receiving real data with selected_real_kind

```
integer :: dptype
real(selected_real_kind(15, 307)) :: x

call MPI_TYPE_CREATE_F90_REAL(15, 307, &
dptype, ierror)

call MPI_Recv(x, 1, dptype, 1, 0, my_comm, &
stat, merror)
```

- Arguments of `selected_real_kind` and `MPI_TYPE_CREATE` need to match exactly!

Fortran kind syntax (F90 syntax) Integer data

```
MPI_TYPE_CREATE_F90_INTEGER(R, NEWTYPE, IERROR)
```

```
INTEGER P, R, NEWTYPE, IERROR
```

- `R`: decimal exponent range
- `NEWTYPE`: requested MPI datatype (handle)
- Use to send integer data declared with `selected_integer_kind`

Fortran kind syntax (F90 syntax)

Complex data

```
MPI_TYPE_CREATE_F90_COMPLEX(P, R, NEWTYPE, IERROR)
```

```
INTEGER P, R, NEWTYPE, IERROR
```

- P: precision decimal digits
- R: exponent range
- NEWTYPE: requested MPI datatype (handle)
- Use to send complex data declared with `selected_real_kind`
- Either P or R may be `MPI_UNDEFINED`

When do MPI calls return

- **MPI_Send**
 - returns: `send-buffer` is safe to be overwritten
 - Whether data has or hasn't arrived depends – no idea!!
- **MPI_Recv**
 - returns:
 - `receive buffer` contains the data
 - you can now use that data
 - in `mpi4py`: the Python object is returned
- Both can involve significant waiting time!!

Timing MPI code

- MPI offers a very nice command to measure performance

- In C:

```
double MPI_Wtime(void)
```

- In Fortran:

```
DOUBLE PRECISION MPI_WTIME()
```

no **IERROR**

- In Python:

```
comm.Wtime()
```

- **MPI_WTIME** returns **seconds** since some time in the past
- Don't use standard Fortran system clock

Summary

- This lecture has introduced
 - Sending point-to-point messages in MPI
 - Most used predefined MPI data types
 - No need to handle datatypes for standard Python objects
- Discussed timing in MPI codes

Exercise π^2

- We have

$$\sum_{n=1}^{\infty} \frac{1}{n^2} = \frac{\pi^2}{6}$$

- Write a serial program to calculate π^2 in C or Fortran
- Parallelise it using MPI
 - When using p processors, each task should do $1/p$ of the number of terms in the sum
 - Each task should work out, the total number of task and the range of terms he needs to work on
 - Once done with his terms, each task sends its partial result to rank 0
 - Rank 0 receives the partial results, calculates the final result and prints it to the screen