

Avoiding deadlock: Non blocking communication


Joachim Hein (LUNARC, Lund University)
Xin Li (PDC, KTH)

Overview

- Eager vs. rendezvous protocol
- Deadlock
- Non-blocking communication
- Performance implications from non-blocking communication

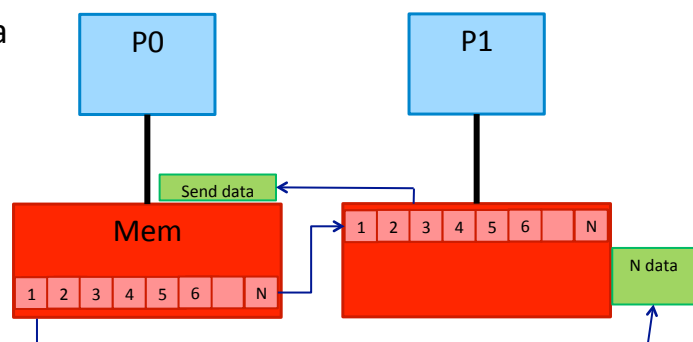
Small messages: Eager protocol

- Eager sending:
- When send is issued, data transferred to buffer on receiver
- There it gets buffered until matching receive is posted
- After receive is posted, copied into process memory
- Eager sending is not suitable for large messages:
- Receiver or Sender: not enough buffer memory to hold transferred data



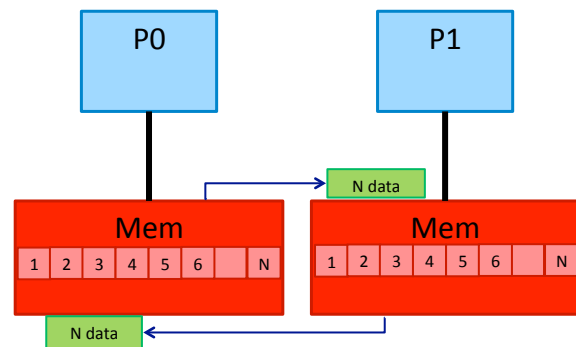
Large messages: Rendezvous protocol

- When send is issued, meta-data transferred, e.g. N integers to come
- Matching receive is posted
 - Memory on receiver available (arg. of rcv)
 - Data request is sent to sender
- Sender sends actual data
 - Sender must not overwrite send buffer until copied
 - Sender will wait forever if no matching receive is posted



Deadlock

- Both tasks use a send operations which goes “rendezvous”
- Both tasks **wait for ever** for the request to send data
- **MPI_Send** can not return, buffer can not be overwritten
- You have to assume that any **MPI_Send** waits for the receive
 - Different system in future
 - Future use more data



Synchronous send

In C:

```
int MPI_Ssend(void* buf, int count, MPI_Datatype
datatype, int dest, int tag, MPI_Comm comm)
```

In Fortran 90:

```
MPI_SSEND(BUF, COUNT, DATATYPE, DEST, TAG, &
COMM, IERROR)
```

In Python:

```
comm.ssend(obj, dest=dest, tag=tag)
```

- Same arguments as MPI_Send
- Call returns only **after** receive has been posted
- Always uses **rendezvous** or similar protocol
- **Useful for testing for deadlock problem in a code**

Non-blocking communication (overview)

- Avoid deadlock by using non-blocking send and/or receive
- Non-blocking calls return directly, in general:
 - Not safe to overwrite a send buffer
 - Receive buffer does not contain data
- Additional calls: **MPI_Wait**, **MPI_Test**, **MPI_Waitall** to check communication has finished
- A “checker” has to be called, otherwise
 - nasty memory or resource leak
 - code may behave **non-deterministic**

No deadlock by non-blocking receive

1. Both tasks issue non-blocking receive
 - Receive buffer becomes available
 - Program continues executing
 2. Both tasks post send “rendezvous” operations
 - Since receive buffers available, send transfers data & returns
 3. Both tasks check on receive (e.g. **MPI_Wait**):
 - Ensures data has arrived in receive buffer
 - Checking operation required – it completes operation
 4. Continue with calculation and use received data
- **Rem:** The non-blocking call continues work, when program is doing other work

Non-blocking receive in C

```
int MPI_Irecv(void* buf, int count, MPI_Datatype
datatype, int source, int tag, MPI_Comm comm,
MPI_Request *request)
```

- buf: address of receive buffer (output)
- count: number of elements to be received
- datatype: data type of buffer (explained further down)
- source: rank of sender (data origin)
- tag: message tag (needs to match the send!)
- comm: communicator
- request: request handle (output), pass into “checker”

Non-blocking receive in Fortran 90

```
MPI_IRECV(BUF, COUNT, DATATYPE, SOURCE, TAG, &
COMM, REQUEST, IERROR)
```

```
<TYPE>:: BUF
```

```
INTEGER:: COUNT, DATATYPE, SOURCE, TAG, &
COMM, REQUEST, IERROR
```

- BUF: address of receive buffer (output)
- COUNT: number of elements to be received
- DATATYPE: data type of buffer (explained further down)
- SOURCE: rank of sender (data origin)
- TAG: message tag (needs to match the send!)
- COMM: communicator
- REQUEST: request handle (output), pass into checker

Non-blocking receive in Python

```
req = comm.irecv(source=source, tag=tag)
```

- `source`: rank of sender (data origin)
- `tag`: message tag (needs to match the send!)
- A `Request` object is returned
- Python object is returned by the `wait` method of `Request` object
- Example:


```
req = comm.irecv(source=0, tag=0)  
obj = req.wait()
```

Non-blocking receive in Python (large data)

```
req = comm.irecv(buf, source=source, tag=tag)
```

- When sending large data via `irecv`, you may see the “message truncated” error.
- In such case a buffer will be needed by `irecv`
- Example:


```
# 1 MB buffer; make it larger if needed  
buf = bytearray(1<<20)  
req = comm.irecv(buf, source=0, tag=0)  
obj = req.wait()
```

Non-blocking send in C: MPI_Isend

```
int MPI_Isend(void* buf, int count, MPI_Datatype
datatype, int dest, int tag, MPI_Comm comm,
MPI_Request *request)
```

- buf: address of send buffer
- count: number of elements to be send
- datatype: data type of buffer (explained further down)
- dest: rank of receiver
- tag: message tag (put 0 if you don't need)
- comm: communicator
- request: request handle (output), pass into checker

Non-blocking send Fortran 90: MPI_ISEND

```
MPI_ISEND(BUF, COUNT, DATATYPE, DEST, TAG, &
COMM, REQUEST, IERROR)
```

- ```
<TYPE>:: BUF
INTEGER:: COUNT, DATATYPE, DEST, TAG, COMM
INTEGER:: REQUEST, IERROR
```
- BUF: address of send buffer
  - COUNT: number of elements to be send
  - DATATYPE: data type of buffer (explained further down)
  - DEST: rank of receiver
  - TAG: message tag (put 0 if you don't need)
  - COMM: communicator
  - REQUEST: request handle (output), pass into checker

## Non-blocking send in Python: `isend`

```
req = comm.isend(obj, dest=dest, tag=tag)
```

- `obj`: the Python object to be sent
- `dest`: rank of receiver
- `tag`: message tag (optional; default is 0)

- A `Request` object is returned
- The `Request` object has the `wait` method

- Example:

```
req = comm.isend(data, source=0, tag=0)
req.wait()
```

## Waiting on non-blocking communication

- In C:

```
int MPI_Wait(MPI_Request *request, MPI_Status *status)
```

- In Fortran 90:

```
MPI_WAIT(REQUEST, STATUS, IERROR)
 INTEGER:: REQUEST, STATUS(MPI_STATUS_SIZE)
 INTEGER:: IERROR
```

- `request`: handle from `MPI_Irecv` or `MPI_Isend`
- `status`: status info from complete operation (recv)

- In Python:

```
req.wait()
```

- Call waits (blocks) until communication has finished



## Example in C

```
int b, a=17;
MPI_Request s_req, r_req;
MPI_Status cstat;

MPI_Irecv(&b, 1, MPI_INT, npr, 0, myc, &r_req);
MPI_Isend(&a, 1, MPI_INT, npr, 0, myc, &s_req);

/* don't touch a or b here, can do other work */

MPI_Wait(&s_req, &cstat);
a = 9; /* now safe to change a */
MPI_Wait(&r_req, &cstat);
b = 2*b; /* b can now be used */
```

## Example in Fortran 90

```
integer:: b, a=17
integer:: s_req, r_req, ierror
integer, dimension(MPI_STATUS_SIZE) :: cstat

MPI_Irecv(b, 1, MPI_INTEGER, npr, 0, myc, &
 r_req, ierror)
MPI_Isend(a, 1, MPI_INTEGER, npr, 0, myc, &
 s_req, ierror)
! don't touch a or b here, can do other work
MPI_Wait(s_req, cstat, ierror)
a = 9 ! now safe to change a
MPI_Wait(r_req, cstat, ierror)
b = 2*b ! b can now be used
```

## Example in Python

```
from mpi4py import MPI

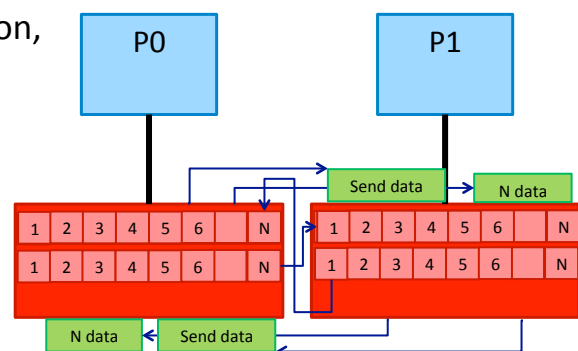
comm = MPI.COMM_WORLD
a = 17

r_req = comm.irecv(source=0, tag=myc)
s_req = comm.isend(a, dest=0, tag=myc)

s_req.wait()
a = 9
b = r_req.wait()
b = 2*b
```

## Example: resolving deadlock: MPI\_Isend

- Both tasks use an **MPI\_Isend** operation
- **Isend** publishes the send operation, e.g. transfers meta data
- **MPI\_Isend** calls return
- Don't modify sendbuffer yet!!
- Both task post **MPI\_Recv**
  - Publish recv-buffers
  - Match of send & recv
  - Request data transfer
  - Receives data
- Tasks must call **MPI\_Wait**
  - Finalises the **Isend** and frees resources



```
int MPI_Test(MPI_Request *request, int *flag,
 MPI_Status *status)
```

**request:** handle from MPI\_Irecv or MPI\_Isend

**flag:** returns **true**, if call has finished

**status:** status info from complete operation (recv)

- In Fortran:

```
MPI_TEST(REQUEST, FLAG, STATUS, IERROR)
```

```
INTEGER :: REQUEST, STATUS(MPI_STATUS_SIZE)
```

```
LOGICAL :: FLAG
```

```
INTEGER :: IERROR
```

- Tests whether communication finished
- De-allocates request handle when true

## Non-blocking communication and performance

- Non blocking communication can be used to **hide communication** time:
  - Issue communication calls early
  - Do other work, not associated with the communication buffers
  - Delay the Wait/Test for the communication calls until data is needed
  - Particular useful on modern RDMA hardware (communication not done by processor)
- Pre-posting non-blocking receives can reduce data traffic
  - Data goes directly into application buffer instead of buffer in MPI library

## Late receive

- If receive not posted, data needs storing somewhere

### Program on P0:


```
a = 13
send(a, 1)
```

### Program on P1:

```
somework(c)
morework(c)

recv(b, 0)

c = c * b
```



## Pre-post receive


- If non-blocking receive pre-posted, can go directly into application space

### Program on P0:

```
a = 13
send(a, 1)
```

### Program on P1:


```
irecv(b, 0, r)
somework(c)
morework(c)
wait(r)
c = c * b
```



## Summary

- Explained the danger of deadlock in p-2-p communication
- Non-blocking communication
  - to avoid deadlock
  - to overlap communication and calculation
  - preposted receives: performance boost on certain architectures
- **Important:** Every non-blocking call (`MPI_Isend`, `MPI_Irecv`, ...) **requires** a matching `MPI_Wait` or a true `MPI_Test`

## Exercise 1



### Ping-Ping code

- Two processes exchange data between them simultaneously
- Each processor should
  - set up a data array of variable length (up to 1 MB of data)
  - send its data to the other processor
  - receive the data from the other processor
  - check received data is correct
- You can use `MPI_Ssend` or `MPI_Isend` to make sure it is not deadlocking
- Time your code using standard and synchronous send

## Exercise

### Messages around a ring

- Each processor should
  - Initialise 2 integer arrays
    - set one to rank number
    - other to 1000x rank number
  - Send one up and down
  - Receive from neighbours
  - Sum up received data in separate (up/down) sums
  - Pass the received data on, to continue in the direction they came
- Checks after final exchange:
  - Last receive (up & down) equal 1x and 1000x own rank
  - The two sums equal:  $\frac{1}{2}n(n-1)$  and  $500n(n-1)$

