

An Introduction to parallel debugging

Joachim Hein (LUNARC)

1

Overview

- Introduction to debugging and parallel debugging
- Running the Linaro DDT parallel debugger

2

INTRODUCTION TO DEBUGGING

3

Traditional standard way to debug: “printf debugging”

- Add extra print statements to the code
 - Indicate whether the code reaches a certain stage
 - Print the values of key variable
- Issues with this approach
 - Need to modify the source code, recompile
 - Iterative approach, frequent recompiles
- Debuggers are more convenient
 - Allows working with unmodified source
 - Allows line by line execution

4

Debuggers

- Linux system come with **gdb** as a debugger
 - Command line execution
 - GUIs exist and are recommended
 - Often integrated into development platforms

5

PARALLEL DEBUGGING

6

Parallel debugging

- Parallel applications offer new levels of complexity
- Before starting, try to simplify the task
 - Problem still there if you **reduce the problem size**?
 - Problem still there if you **reduce the task/thread count**?
- “printf debugging” even more problematic than in serial
 - More output (different tasks/threads printing)
 - Identification of task/thread printing required
 - UNIX grep helpful to filter output

7

Parallel debuggers

- Licenses are expensive
 - Being able to do “printf” is an essential skill
- Parallel debuggers became more usable over the years
- I am aware of two products
 - Totalview for HPC (<https://totalview.io/>)
 - Linaro DDT - part of Linaro Forge
 - Formerly known as ALLINEA DDT/FORGE and ARM DDT/FORGE
 - There is a NAISS wide license

8

PREPARATIONS AND STARTING DDT

9

Preparations

- HPC system needs to display the gui on your monitor
 - VNC solution (e.g. LUNARC HPC desktop, ThinLinc)
 - Connect with X-forwarding (ssh -X ...)
- Recompile your application with the flags: -g -O0


```
mpif90 -g -O0 -o hello_mpi hello_mpi.f90
```
- Comments:
 - Lack of optimisation slows code (in particular C++)
 - Problem might disappear – hint for overrun array
 - You can use optimisation
 - Though match code line to instruction might not work

10

Start the gui

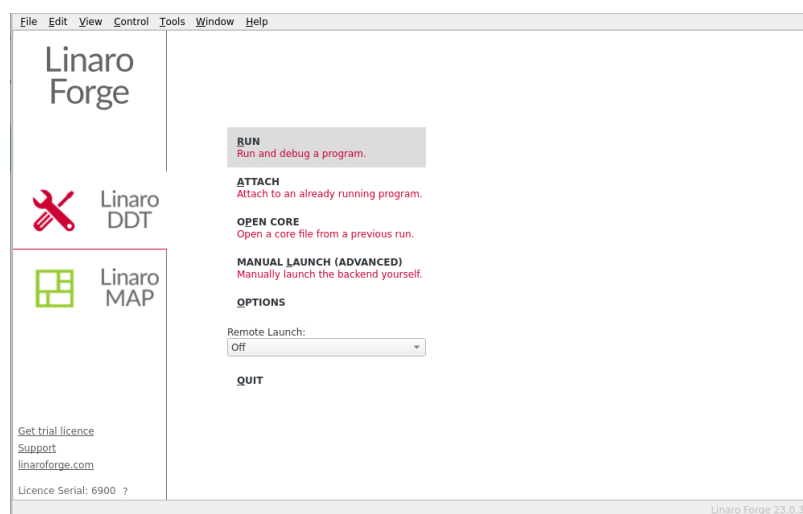
- Best to start the gui on the login node and keep it running

```
module load linaro_forge/23.0.3
ddt &
```

- Alternative use the ARM remote client on your desktop and connect to the (front end of the)HPC service

11

Linaro Forge gui



12

Starting code on the compute nodes

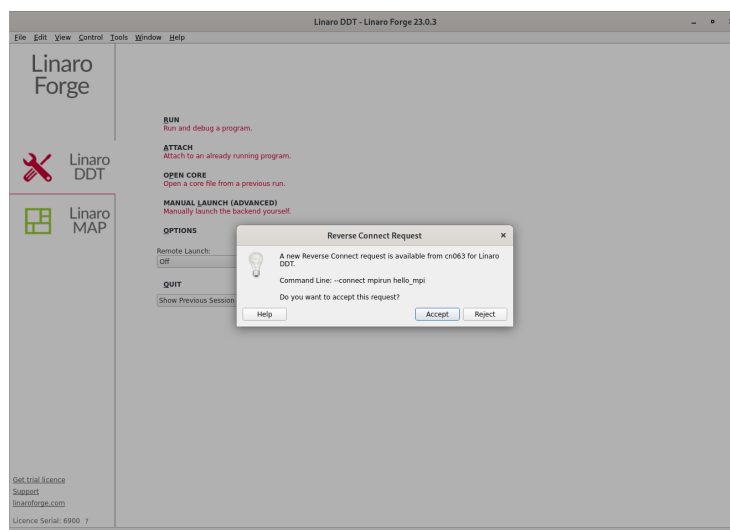
- Transfer to the backend node
 - Jobscript
 - Interactive allocation
- Make sure relevant modules are loaded
 - compiler, MPI lib, other libs, Linaro DDT/Forge
- Prefix job launcher with: `ddt --connect`

```
ddt --connect mpirun mpihello
```

```
ddt --connect mpirun python3 %allinea_python_debug% hello_mpi.py
```

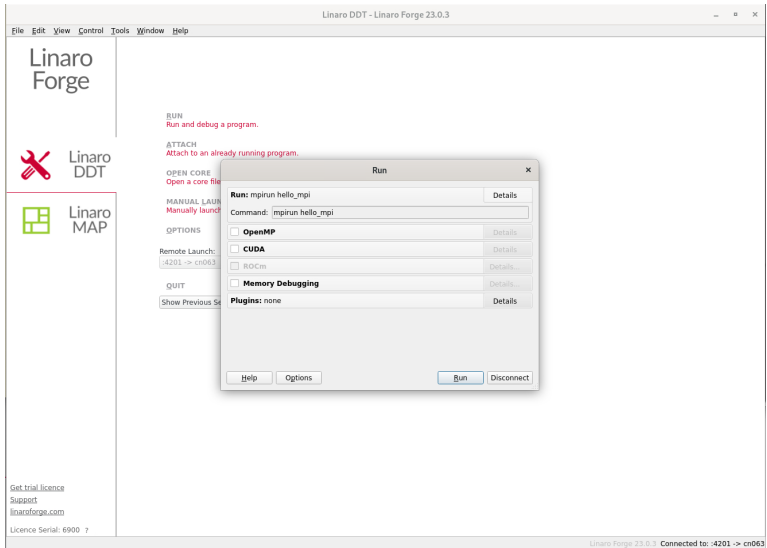
13

Accept the “Reverse Connect request”



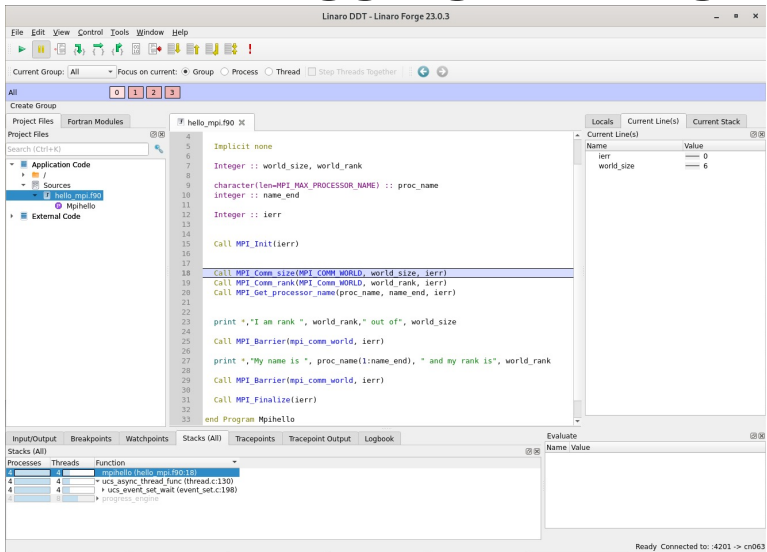
14

Start running your program



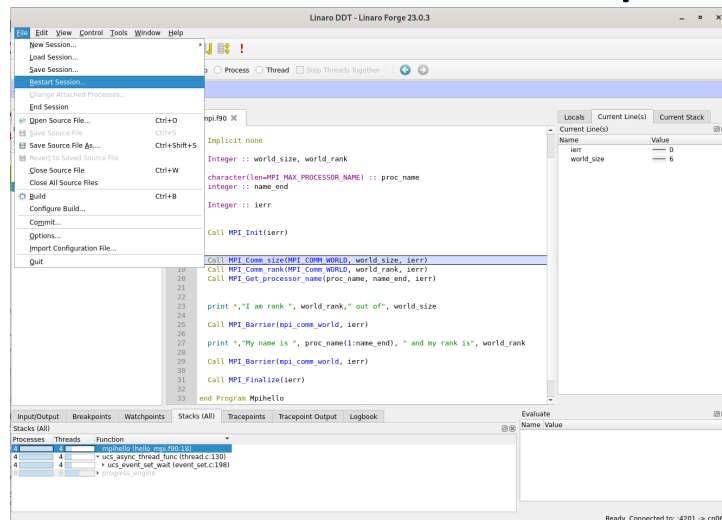
15

Start debugging in the gui



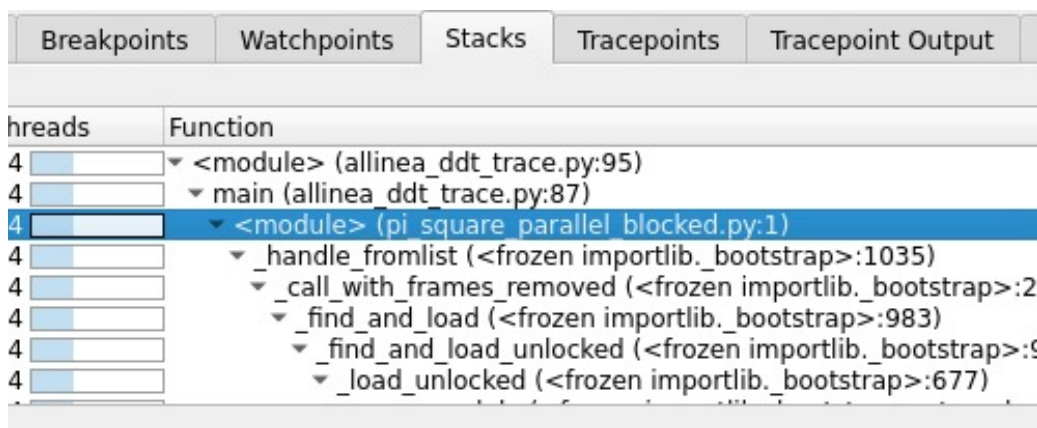
16

Starting over – frequently required Use “Restart session” option



17

Starting a Python debug session



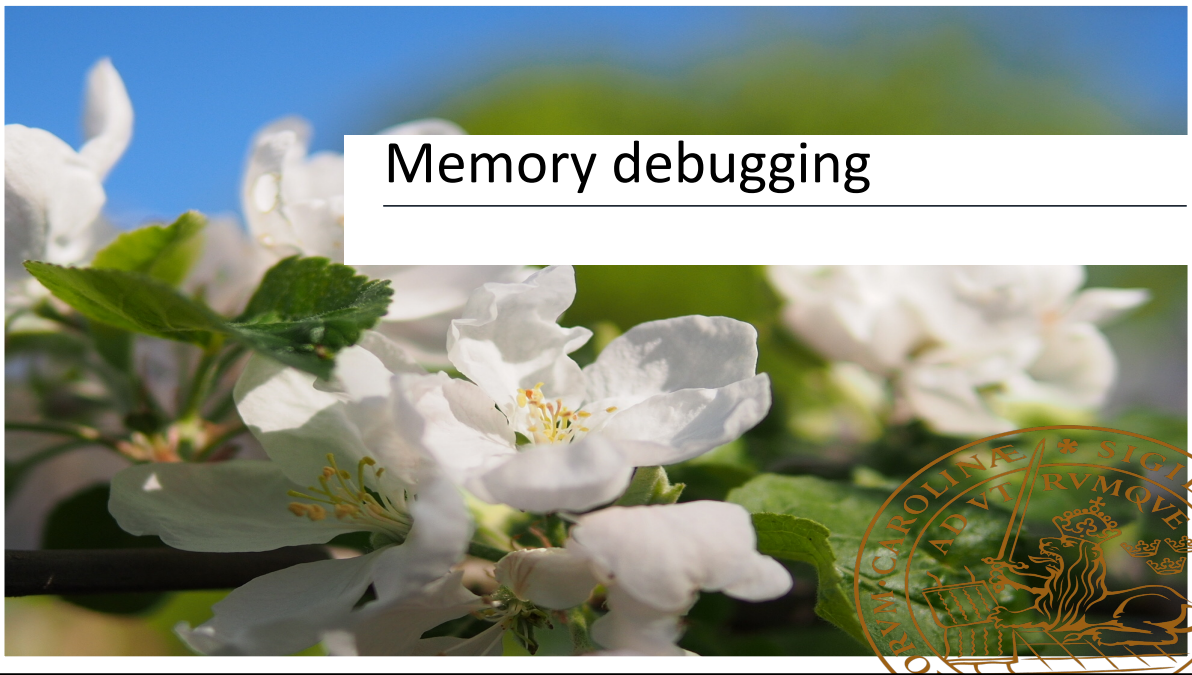
- You might need to locate Python script in “Stacks” window
- Set breakpoint on first line after importing (from) mpi4py

18

Demo

- hello world (Fortran)
- Message on a ring (C)
- Pi-square (Python)

19



Memory debugging

20

Problematic memory access

- Codes often suffer from memory problems
 - Writing in memory locations they shouldn't
 - Illegal deallocation (double, bad pointer position, ...)
 - Memory leaks
- Typical signatures of memory problems
 - Seg-faults
 - Code behaviour changes when:
 - Editing (e.g. printf debugging)
 - Changing compilers or optimisation flags

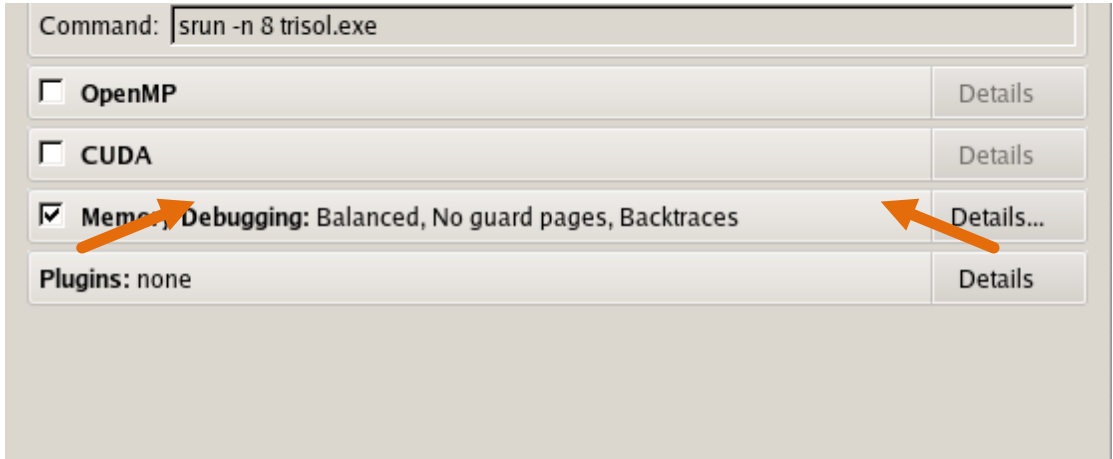
21

Activating memory debugging in DDT

- Replace the malloc library with ARM's dmalloc
- Comes in 4 versions:
 - C/Fortran no threads
 - C/Fortran threads
 - C++ no threads
 - C++ threads
- Current version seems to prefer “threads”

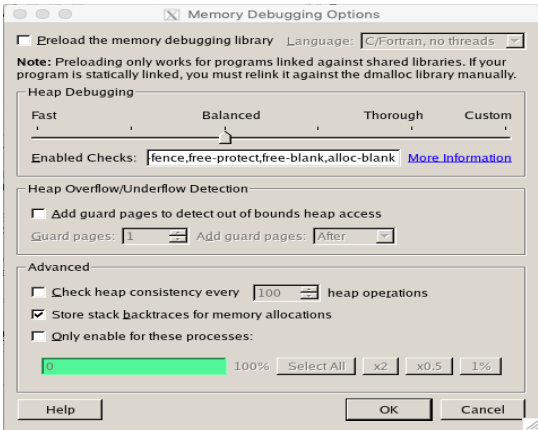
22

Select memory debugging



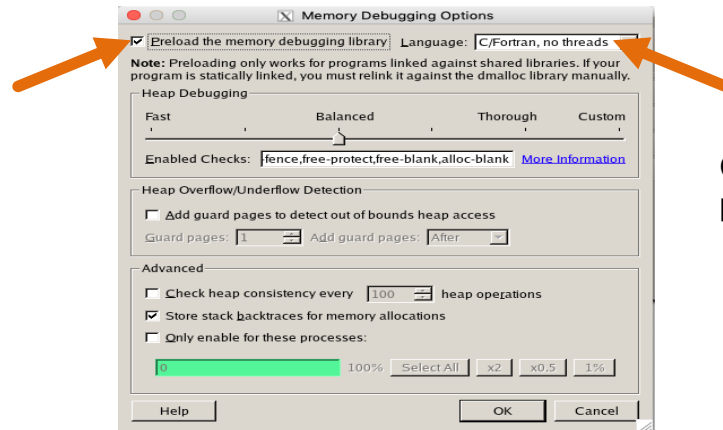
23

Selecting Memory Debugging Option



24

Dynamic linking



Current version
prefers: threads!

25

Static linking

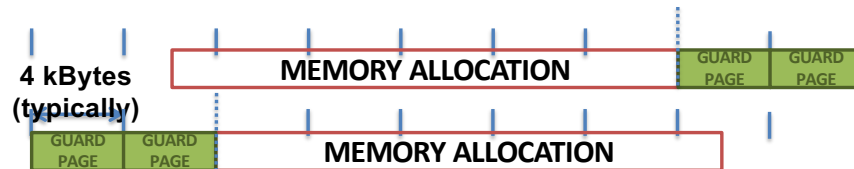
- If you link statically or if dynamic linking fails
- Add a line like (check user guide)

```
-Wl,--allow-multiple-definition,--undefined=malloc /path/lib/64/libdmalloc.a
```

to the link line **before** anything else
– Often required on CRAYs

26

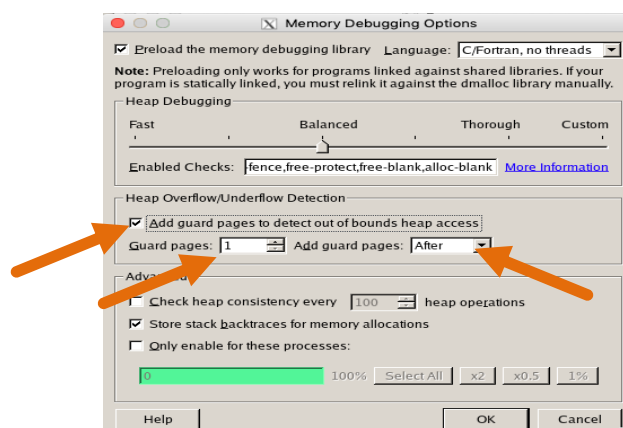
Guard pages (aka “electric fences”)



- **A powerful feature...:**
 - Forbids read/write on guard pages throughout the whole execution
(because it overrides C Standard Memory Management library)
- **... to be used carefully:**
 - Kernel limitation: up to 32k guard pages max
 - Beware the additional memory usage cost

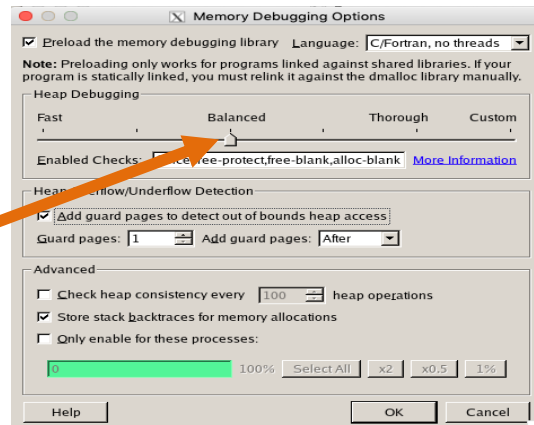
27

Activate guard pages



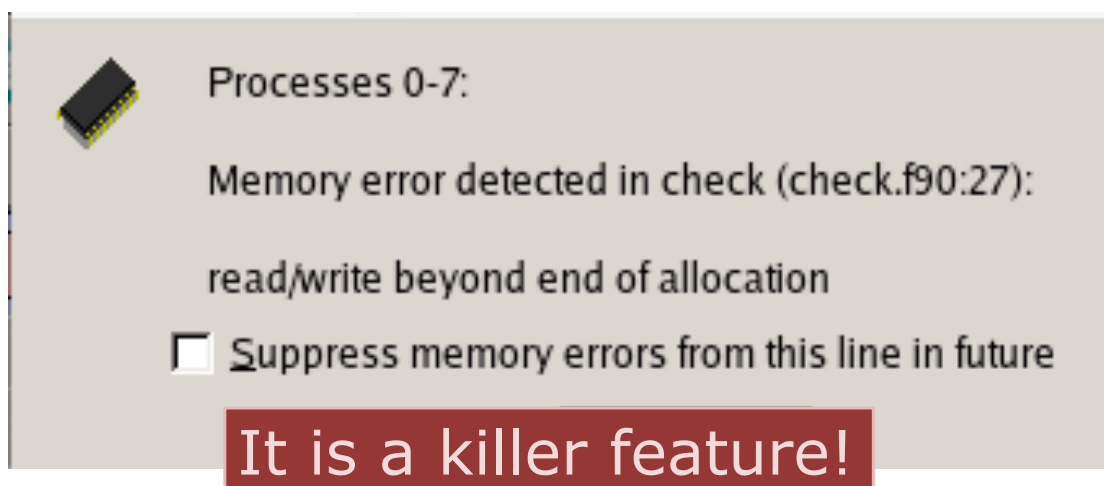
28

Select the depth of the tests



29

When it finds something you get:



30

Demo

- Locating memory issue

31

Recap/Summary

- Starting the gui
- Demonstrating how to run it
- Memory debugging feature
 - This saved me so much time in the years

32

Acknowledgements

- Juan Gao (then at ARM)
- Patrick Wohlschlegel (than at ARM)
- Thor Wikfeldt (than at KTH/PDC)