

# MPI Performance

Pedro Ojeda, Joachim Hein

HPC2N, Umeå University

LUNARC & Centre of Mathematical Sciences

Lund University

# Speedup and efficiency

Let us call  $T_{\text{serial}}$  the time a program takes by using a single core and  $T_{\text{parallel}}$  the time it takes by using  $p$  processors, the speedup is defined by:

$$S = \frac{T_{\text{serial}}}{T_{\text{parallel}}}$$

And the efficiency by:

$$E = \frac{S}{p} = \frac{T_{\text{serial}}}{p * T_{\text{parallel}}}$$

Can we achieve  $T_{\text{parallel}} = T_{\text{serial}}/p$ ?

# Speedup, efficiency and cost

Let us call  $T_{\text{serial}}$  the time a program takes by using a single core and  $T_{\text{parallel}}$  the time it takes by using  $p$  processors, the speedup is defined by:

$$S = \frac{T_{\text{serial}}}{T_{\text{parallel}}}$$

And the efficiency by:

$$E = \frac{S}{p} = \frac{T_{\text{serial}}}{p * T_{\text{parallel}}}$$

The cost is:

$$C = p * T_{\text{parallel}}$$

In general,  $T_{\text{parallel}} = \frac{T_{\text{serial}}}{p} + T_{\text{overhead}}$

# Amdahl's law

Let us call the proportion of the code that can be parallelized by  $\tau$ , the parallel time using  $p$  processors can be written as,

$$T_{\text{parallel}} = \tau * \frac{T_{\text{serial}}}{p} + (1 - \tau) * T_{\text{serial}}$$

Then, the speedup will be

$$S = \frac{T_{\text{serial}}}{T_{\text{parallel}}} = \frac{1}{1 - \tau * \left(1 - \frac{1}{p}\right)}$$

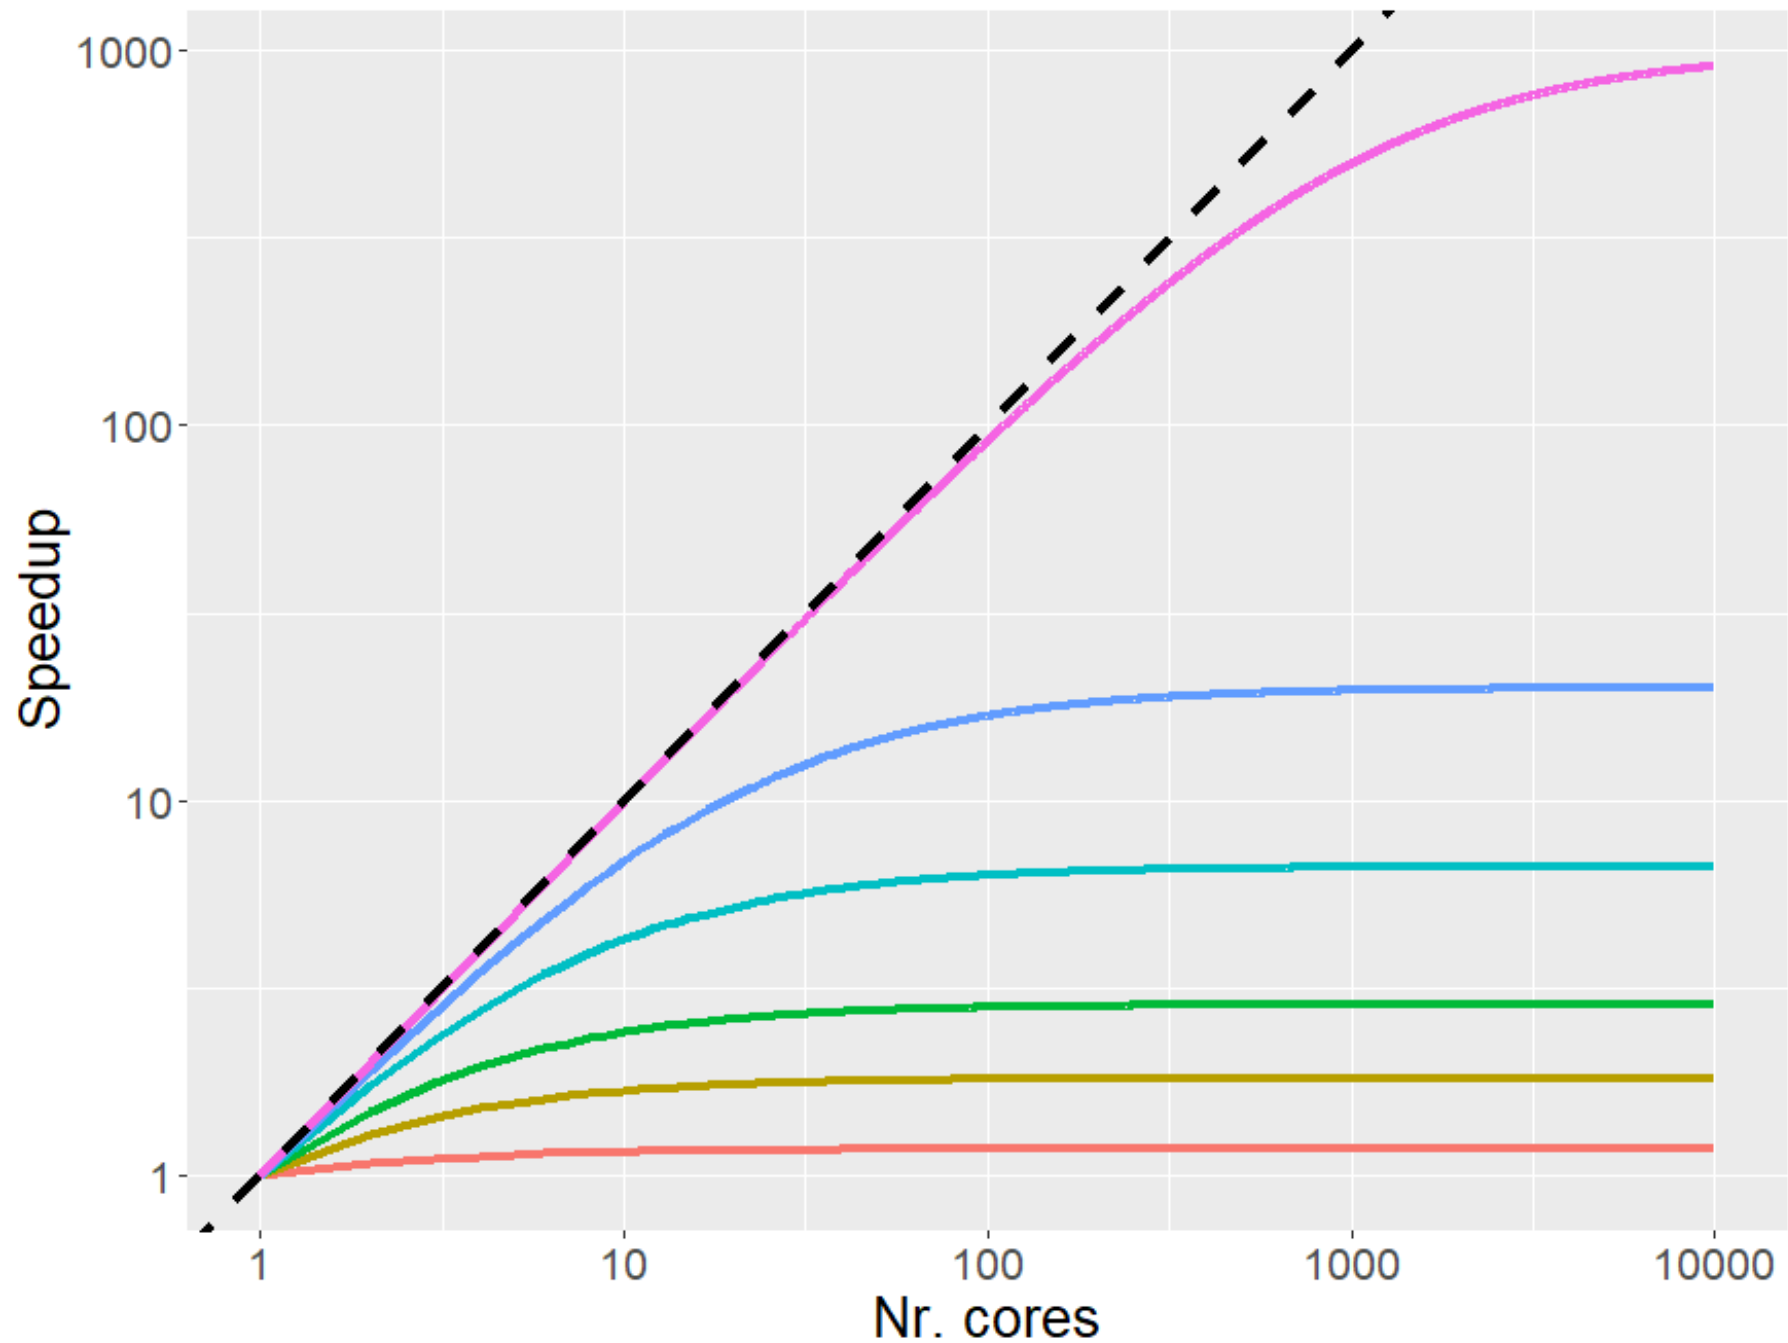
# Amdahl's law

Let us call the proportion of the code that can be parallelized by  $\tau$ , the parallel time using  $p$  processors can be written as,


$$T_{\text{parallel}} = \tau * \frac{T_{\text{serial}}}{p} + (1 - \tau) * T_{\text{serial}}$$

Then, the speedup will be

$$S = \frac{T_{\text{serial}}}{T_{\text{parallel}}} = \frac{1}{1 - \tau * \left(1 - \frac{1}{p}\right)}$$



$$\lim_{p \rightarrow \infty} S = \frac{1}{1 - \tau}$$



This behavior is known as **strong scaling**

# Speedup in practice

For a certain code the speedup is computed as follows:

$$S = \frac{T_{\text{serial}}}{T_{\text{parallel}}}$$

Where  $T_{\text{serial}}$  is the simulation time with a single core and  $T_{\text{parallel}}$  is the time obtained by using 1,2,...p processors.

# Gustafson's law

In general, the goal of parallelizing a code is not only to use a large number of cores but also to simulate larger problems in size.

J. L. Gustafson reevaluated the Amdahl's law<sup>1</sup> for problems with a variable size and gave an alternative based on E. Barsis suggestion:

$$\text{Scaled speedup} = 1 + \tau * (p - 1)$$

where  $\tau$  and  $p$  have the same meaning as before. Scaled speedup has a linear dependence with the number of processors. This is known as the **weak scaling** behavior.

<sup>1</sup>Communications of the ACM, J. L. Gustafson (1988)



# Scaled speedup in practice

For a certain code the Scaled speedup is computed as follows:

$$\text{Scaled speedup} = \frac{p * T_{\text{serial}}}{T_{\text{parallel}}}$$

Where  $T_{\text{serial}}$  is the simulation time with a single core and  $T_{\text{parallel}}$  is the time obtained by using 1,2,...p processors while changing the size of the simulated problem.

# Timing in MPI

The function `MPI_Wtime()` allows you to set a reference point in time, a time interval can be obtained by taking the difference between two calls of this function.

In C:

```
double MPI_Wtime()
```

In Fortran 90:

```
double precision MPI_WTIME()
```

The time resolution of the hardware can be obtained with `MPI_Wtick()`

# Timing in MPI

Example in Fortran 90:

```
double precision begin_t, end_t  
  
begin_t = MPI_WTIME()  
  
... code to be monitored  
  
end_t = MPI_WTIME()  
  
print *, 'time ', end_t - begin_t, 'sec.'
```

Example in C:

```
double begin_t, end_t;  
  
begin_t = MPI_Wtime();  
  
... code to be monitored  
  
end_t = MPI_Wtime();  
  
printf("time %1.2f sec.", end_t - begin_t);
```

# Choice of algorithms


Consider the problem of computing the matrix vector multiplication  $\mathbf{M}_{n \times n} \mathbf{b}_n = \mathbf{c}_n$  using a parallel algorithm. The total number of operations is  $2n^2 - n$  and the total time for computation is  $(2n^2 - n) * T_{10p}$  ( $T_{10p}$  being the time for a single operation).

If each row is transferred to the processors (for each element of  $\mathbf{c}$ ) and the result is transferred back to the master, the total time spent in data movement is  $(n^2 + n) * T_{comm}$  ( $T_{comm}$  being the time for transfer a single element). Thus, the ratio of computation and communication is

$$r = \left( \frac{n^2 + n}{2n^2 - n} \right) * \left( \frac{T_{comm}}{T_{10p}} \right) \quad \text{and} \quad \lim_{n \rightarrow \infty} r \propto \frac{1}{2}$$


# Overhead in MPI programs

$$T_o = p * T_{\text{parallel}} - T_{\text{serial}}$$



MPI profile from a GROMACS simulation obtained with Extrae/Paraver showing the MPI calls metric in *Profiling and Tracing Tools for Performance Analysis of Large Scale Applications*, J. Eriksson, et. al. (<https://prace-ri.eu/wp-content/uploads/WP237.pdf>)


# Typical MPI Point-to-Point performance: Aurora



# MPI Point-to-Point performance

- Typical latency (minimum time to send anything):
  - between 1 and 5  $\mu\text{sec}$  via network
  - around 1  $\mu\text{sec}$  inside node (shared memory/cache)
  - these are  $O(10000 \text{ cycles})$  - an eternity!
- Modern hardware transfers several GB per second
  - min message size for that  $\approx 10 \text{ kB}$
- For message smaller than  $\approx 1 \text{ kB}$ 
  - time (almost) independent of message size
- Aim to **aggregate your messages**
  - Always try to send more than a few numbers
  - Still don't bloat

# Load imbalance



Source: *GROMACS: High performance molecular simulations through multi-level parallelism from laptops to supercomputers*, M. J. Abraham, *SoftwareX*, 1, 19, (2015)

Fig. 2. *Left*: The protein lysozyme (24,119 atoms) in a compact unit cell representation corresponding to a rhombic dodecahedron with close-to-spherical shape. *Right*: Internally, this is represented as a triclinic cell and a load-balanced staggered  $6 \times 4 \times 3$  domain decomposition grid on 72 MPI ranks. The PME lattice sum is calculated on a uniform grid of  $6 \times 4$  MPI ranks (not shown).



# Summary

- Some performance metrics were studied in this section: Speedup, Efficiency, Cost, Overhead, and Scaled speedup. Depending on the parallelization goals, one of them could be more appropriate than the others.
- Strong and weak scaling concepts were introduced.
- The choice of the parallel algorithm is of major importance.
- Regarding data transfer, always aim to send fewer but larger messages

# Exercise

- Use the provided code for the 2D integration to get a plot of the Speedup (keep the number of bins constant and change the number of processors)

(Optional) Use the Amdahl's law to fit a curve to your results and get the fraction of the parallel/serial time

- Obtain a plot of the Scaled speedup for the same code (change both the number of bins and the number of processors)

(Optional) Use the Gustafson's law to fit a curve to your results and get the fraction of the parallel/serial time